



**dd**

# What is dd?



A backup utility

Performs a sequential, byte-for-byte copy

## What can dd work with?

dd can use any file on the system

# In Unix, everything is a file



This includes (but not limited to):

Files

Folders

**Devices**

Processes

“dd can use any file” includes any **device** on the system

# Example 1



Backing up a hard drive to an image:

```
dd if=/dev/sda of=disk.img
```

First hard drive

Stands for "input file"

Disk image file

Stands for "output file"

This works with any disk, including CDs.

# Example 2



Restoring an image to a hard drive:

```
dd if=disk.img of=/dev/sda
```

All I did was reverse input and output

We can use the mount command to open this image for reading/writing

# Example 3



Compressing a disk image (while making it):

```
dd if=/dev/sda | gzip -c > disk.img.gz
```

Standard output from dd is taken to the standard input to gzip (which is then written to a file)

This also works in reverse:

```
gunzip -c disk.img.gz | dd of=/dev/sda
```

# Example 4



Formatting and/or writing zeros to a disk will erase data, however there is the possibility of recovery

We need to write random 1's and 0's to ensure data cannot be recovered

```
dd if=/dev/random of=/dev/sda
```

Using Linux's random number generator, this will securely erase the entire disk

# Example 5



What if the partition table is corrupt, and you need to restore partitions?

If you know where the partitions are, you can clear out the table, then restore the partition

```
dd if=/dev/zero of=/dev/sda bs=1 count=512
```

This will clear the partition table, and allow you to use a utility (such as fdisk) to restore the partitions



# Conclusion



dd is a simple tool with a simple job

Thanks to Linux, we can use it for a variety of tasks

If we tie dd together with other simple programs and utilities, we find a wealth of functionality hidden in a few simple tools.

# Appendix A: more backup



To properly backup a partition:

```
dd if=/dev/sda of=output.img bs=4k  
conv=noerror,notrunc
```

`bs=4k`

Means “read 4 blocks at a time”

Ensures we properly read blocks from the filesystem

`conv=noerror`

Means “do not stop for errors”

# Appendix A (cont.)



`conv=notrunc`

Means “do not truncate output”

We use this to prevent dd from padding the output with extraneous zeros

# Appendix B: mount



If we made a disk image of a single partition, we can use mount to load it into an empty folder just like it was a hard disk

```
mount -t <fs> -o loop partition.image mount/
```

Loop tells Linux to treat the image as a device

If we did an entire disk, we need to give mount an *offset*

This offset is set to the physical position of the partition we want in the disk.

# Appendix B (cont.)



If we have a disk image of several partitions,  
mount will need an offset to mount one of them

fdisk -l will show where the partitions begin  
(marked by 'start')

fdisk lists things in blocks, but mount's offset is in  
bytes

512 bytes per block, so multiply the start number  
by that to get the partition's offset

The final command

```
mount -t <fs> -o loop,<offset> image.img mount/
```

# Appendix C: using bs



Normally, we specify `bs=4k` (or 4096)

Different hardware will have different optimal block sizes, but we can easily test that

```
dd if=/dev/zero of=<destination> count=10  
bs=<blocksize>
```

This will count 10 blocks, then stop and report the transfer speed

You can use this to compare the write speed of different block sizes



**dd**

# What is dd?



A backup utility

Performs a sequential, byte-for-byte copy

## What can dd work with?

dd can use any file on the system



# In Unix, everything is a file



This includes (but not limited to):

Files

Folders

**Devices**

Processes

“dd can use any file” includes any **device** on the system

## Example 1



Backing up a hard drive to an image:

```
dd if=/dev/sda of=disk.img
```

Stands for "input file"      First hard drive      Disk image file

Stands for "output file"

This works with any disk, including CDs.

## Example 2



Restoring an image to a hard drive:

```
dd if=disk.img of=/dev/sda
```

All I did was reverse input and output

We can use the `mount` command to open this image for reading/writing

## Example 3



Compressing a disk image (while making it):

```
dd if=/dev/sda | gzip -c > disk.img.gz
```

Standard output from dd is taken to the standard input to gzip (which is then written to a file)

This also works in reverse:

```
gunzip -c disk.img.gz | dd of=/dev/sda
```

## Example 4



Formatting and/or writing zeros to a disk will erase data, however there is the possibility of recovery

We need to write random 1's and 0's to ensure data cannot be recovered

```
dd if=/dev/random of=/dev/sda
```

Using Linux's random number generator, this will securely erase the entire disk

## Example 5



What if the partition table is corrupt, and you need to restore partitions?

If you know where the partitions are, you can clear out the table, then restore the partition

```
dd if=/dev/zero of=/dev/sda bs=1 count=512
```

This will clear the partition table, and allow you to use a utility (such as fdisk) to restore the partitions

# Conclusion



dd is a simple tool with a simple job

Thanks to Linux, we can use it for a variety of tasks

If we tie dd together with other simple programs and utilities, we find a wealth of functionality hidden in a few simple tools.

## Appendix A: more backup



To properly backup a partition:

```
dd if=/dev/sda of=output.img bs=4k  
conv=noerror,notrunc
```

**bs=4k**

Means “read 4 blocks at a time”

Ensures we properly read blocks from the filesystem

**conv=noerror**

Means “do not stop for errors”



## Appendix A (cont.)



`conv=notrunc`

Means “do not truncate output”

We use this to prevent dd from padding the output with extraneous zeros

## Appendix B: mount



If we made a disk image of a single partition, we can use mount to load it into an empty folder just like it was a hard disk

```
mount -t <fs> -o loop partition.image mount/
```

Loop tells Linux to treat the image as a device

If we did an entire disk, we need to give mount an *offset*

This offset is set to the physical position of the partition we want in the disk.

## Appendix B (cont.)



If we have a disk image of several partitions,  
mount will need an offset to mount one of them

fdisk -l will show where the partitions begin  
(marked by 'start')

fdisk lists things in blocks, but mount's offset is in  
bytes

512 bytes per block, so multiply the start number  
by that to get the partition's offset

The final command

```
mount -t <fs> -o loop,<offset> image.img mount/
```

## Appendix C: using bs



Normally, we specify `bs=4k` (or `4096`)

Different hardware will have different optimal block sizes, but we can easily test that

```
dd if=/dev/zero of=<destination> count=10  
    bs=<blocksize>
```

This will count 10 blocks, then stop and report the transfer speed

You can use this to compare the write speed of different block sizes