

# Linux File Systems Part 1

Panther Linux User Group

By: Michael Tracy & Alwyn Malachi Berkeley


# What is a File System?

- What is a File? Abstraction consisting of 0s and 1s
- A method for storing, manipulating, and retrieval computer files
- Think of it as a database for storage mediums

# How a Unix FS works

- Files are represented by inodes
- Inodes denote the blocks allocated to a file
- A folder is simply a reference to inode numbers
- Files are just inodes

# Diagram



Taken from Wikipedia ([http://en.wikipedia.org/wiki/Hard\\_disk\\_drive](http://en.wikipedia.org/wiki/Hard_disk_drive))

# Now We Can Talk About FS

- Organize clusters/blocks
- Saves file information in FAT & inode
- Only the basics
- You can secure basic operations through ACL or Capabilities
- Different types for different media
- Journaling, versioning, log-structured, database, transactional, network, special purpose



# Ext2

- 'The Second Extended filesystem'
- Successor to Ext1 and MinixFS
- Features:
  - Files are referenced by inodes
  - No journal
  - Rock solid—FS of choice for small installations

# Ext3

- 'Third extended filesystem'
- Features:
  - Journal
  - Larger file limits than ext2
- Criticism
  - No online defrag
  - No dynamic inode allocation

# Ext4

- 'Fourth extended filesystem'
- Features:
  - Larger limits and better performance over Ext3
  - Extents
- Criticism:
  - Write through hole



# ReiserFS 3

Circular log

1. remove entry, delete data

Hans Reiser

Files under 4 kilobytes

Defragmentation

Online resizing, growth only

Big Kernel Lock, scaling on multicore

Fsck could destroy internal files

# ReiserFS 4

- Defragmenting ability
- Tailpacking, hit performance
  - Encryption support
- Faster handling of directories /w many files

# New Filesystems

# Tux 3

- Daniel Phillip
- Zumastor NAS project
- Snapshots & Remote Replication

# Server Filesystems

# XFS/JFS

- The original journaling filesystems for Linux
- Very fast for large files and very large partitions
- Features:
  - Extents
  - Dynamic allocation
  - Striping of files



ZFS

See Wikipedia...

# BtrFS (*“Butter FS”*)

Dynamic inode Allocation

Online utilities for resizing etc.

Snapshots

Many similarities to zfs

# Flash File Systems

- Disk file systems are not optimal on flash drives (Erase blocks first, random access(seeks), writes distributed)



JFFS2

UNDER CONSTRUCTION

# SquashFS

Embedded systems

# Network Filesystems



# SMBFS/CIFS

- Windows Shares

# ProcFS

- If the kernel was a person, this is what the person would be thinking